

STATS 506 Problem Set 6

Lingzhi Hao

GitHub Repository:

<https://github.com/Lingzhi-Hao/STATS-506-Problem-Set-06>

Problem 1 - Rcpp

```
library(Rcpp)
```

Warning: package 'Rcpp' was built under R version 4.5.2

```
sourceCpp(code = '
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]
double C_moment(NumericVector x, int k) {
  int n = x.size();
  if (n == 0) {
    stop("x has length 0");
  }

  double mu = 0.0;
  for (int i = 0; i < n; ++i) mu += x[i];
  mu /= n;

  double acc = 0.0;
  for (int i = 0; i < n; ++i) acc += std::pow(x[i] - mu, k);

  return acc / n;
}
')
```

```
set.seed(123)
x <- rnorm(1000)
k <- 3

C_moment(x, k)
```

```
[1] 0.06358509
```

```
e1071::moment(x, order = k, center = TRUE)
```

```
[1] 0.06358509
```

Problem 2 - Expanding on waldCI

(a)

```
source("waldCI.R")
```

```
setClass("bootstrapWaldCI",
        contains = "waldCI",
        slots = c(
            data = "ANY",
            statistic_fn = "function",
            reps = "numeric"
        )
)

##' Compute bootstrap Wald CI
##' @param statistic_fn Function to compute the scalar statistic of interest
##' @param data Data set to use for bootstrap
##' @param reps Number of bootstrap replicates
##' @param level Confidence level (default 0.95)
##' @param compute Computation method: "serial" or "parallel" (default "serial")
##' @return A `bootstrapWaldCI` object
##' @export
makeBootstrapCI <- function(statistic_fn,
                               data,
                               reps,
                               level = 0.95,
```

```

            compute = "serial") {

stopifnot(is.function(statistic_fn))
stopifnot(is.numeric(reps) && length(reps) == 1 && reps > 0)
stopifnot(compute %in% c("serial", "parallel"))

boot_result <- .runBootstrap(statistic_fn, data, reps, compute)

boot_mean <- mean(boot_result)
boot_sterr <- sd(boot_result)

new("bootstrapWaldCI",
  level = level,
  mean = boot_mean,
  sterr = boot_sterr,
  data = data,
  statistic_fn = statistic_fn,
  reps = reps)
}

.runBootstrap <- function(statistic_fn, data, reps, compute) {
  n <- nrow(data)

  if (compute == "serial") {
    results <- numeric(reps)
    for (i in 1:reps) {
      indices <- sample(1:n, replace = TRUE)
      boot_data <- data[indices, , drop = FALSE]
      results[i] <- statistic_fn(boot_data)
    }
    return(results)
  } else if (compute == "parallel") {
    library(parallel)

    num_cores <- detectCores() - 1
    if (num_cores < 1) num_cores <- 1
    cl <- makeCluster(num_cores, type = "PSOCK")

    clusterExport(cl, c("statistic_fn", "data", "n"), envir = environment())
  }
}

```

```

boot_run_fn <- function(i) {
  indices <- sample(1:n, replace = TRUE)
  boot_data <- data[indices, , drop = FALSE]
  statistic_fn(boot_data)
}

results <- parLapply(cl, 1:reps, boot_run_fn)
stopCluster(cl)

return(unlist(results))
}
}

####' Perform a new bootstrap on an existing bootstrapWaldCI object
####' @param object A `bootstrapWaldCI` object
####' @return A new `bootstrapWaldCI` object
####' @export
rebootstrap <- function(object) {
  stopifnot(is(object, "bootstrapWaldCI"))

  boot_result <- .runBootstrap(object@statistic_fn,
                                data = object@data,
                                reps = object@reps,
                                compute = "serial")
  boot_mean <- mean(boot_result)
  boot_sterr <- sd(boot_result)

  return(new("bootstrapWaldCI",
            level = object@level,
            mean = boot_mean,
            sterr = boot_sterr,
            data = object@data,
            statistic_fn = object@statistic_fn,
            reps = object@reps))
}

```

(b)

```
ci1 <- makeBootstrapCI(function(x) mean(x$y),
                        ggplot2::diamonds,
                        reps = 1000)
ci1
```

95% CI: (5.724897, 5.744366)

```
rebootstrap(ci1)
```

95% CI: (5.725073, 5.744577)

```
set.seed(123)

time_serial <- system.time(
  ci1_serial <- makeBootstrapCI(
    function(x) mean(x$y),
    ggplot2::diamonds,
    reps      = 1000,
    compute   = "serial"
  )
)

cat("Serial Computation Results\n")
```

Serial Computation Results

```
cat("Time:\n")
```

Time:

```
print(time_serial)
```

user	system	elapsed
5.31	0.75	8.01

```
cat("CI:\n")
```

CI:

```
ci1_serial
```

95% CI: (5.724917, 5.744346)

```
rebootstrap(ci1_serial)
```

95% CI: (5.725063, 5.744586)

```
set.seed(123)
```

```
time_parallel <- system.time(
  ci1_parallel <- makeBootstrapCI(
    function(x) mean(x$y),
    ggplot2::diamonds,
    reps      = 1000,
    compute   = "parallel"
  )
)
```

```
cat("Parallel Computation Results\n")
```

Parallel Computation Results

```
cat("Time:\n")
```

Time:

```
print(time_parallel)
```

user	system	elapsed
0.05	0.07	14.15

```
cat("CI:\n")
```

CI:

```
ci1_parallel
```

95% CI: (5.725323, 5.744062)

```
rebootstrap(ci1_parallel)
```

95% CI: (5.724917, 5.744346)

Using 1000 bootstrap replications, the parallel computation (23.44s) was slower than the serial computation (17.85s).

Parallel computation in R has substantial overhead from creating worker processes, exporting data, and collecting results. When each bootstrap task is lightweight and the number of replications is moderate, the overhead dominates and parallel method does not run faster. If we have larger reps or more computationally expensive tasks, the parallel method would likely perform better than the serial method.

(c)

```
##' Fits the model and returns the coefficient for disp
##' @param data A data frame
##' @return The estimated coefficient for disp
dispCoef <- function(data) {
  fit <- lm(mpg ~ cyl + disp + wt, data = data)
  return(coef(fit)["disp"])
}
```

```
ci2 <- makeBootstrapCI(dispCoef,
                        mtcars,
                        reps = 1000)
ci2
```

95% CI: (-0.01005764, 0.0250501)

```
rebootstrap(ci2)
```

95% CI: (-0.0114781, 0.02544108)

```
set.seed(123)
```

```
cat("Serial Computation Results\n")
```

Serial Computation Results

```
time_serial_c <- system.time(
  ci2_serial <- makeBootstrapCI(
    dispCoef,
    mtcars,
    reps      = 1000,
    compute   = "serial"
  )
)
```

```
cat("Serial Compute Time (seconds):", time_serial_c["elapsed"], "\n")
```

Serial Compute Time (seconds): 0.89

```
cat("CI:\n")
```

CI:

```
ci2_serial
```

95% CI: (-0.01111183, 0.02598235)

```
rebootstrap(ci2_serial)
```

95% CI: (-0.01230038, 0.02644671)

```
cat("\nParallel Computation Results\n")
```

Parallel Computation Results

```

set.seed(123)

time_parallel_c <- system.time(
  ci2_parallel <- makeBootstrapCI(
    dispCoef,
    mtcars,
    reps      = 1000,
    compute   = "parallel"
  )
)

cat("Parallel Compute Time (seconds):", time_parallel_c["elapsed"], "\n")

```

Parallel Compute Time (seconds): 0.71

```
cat("CI:\n")
```

CI:

```
ci2_parallel
```

95% CI: (-0.01113828, 0.02562702)

```
rebootstrap(ci2_parallel)
```

95% CI: (-0.01111183, 0.02598235)

Both the serial method (2.59s) and parallel method (3.53s) are very fast, but the parallel method is slower.

Similar to (b), parallel computation has overhead from creating worker processes, exporting data, and collecting results. Since computing the coefficient of a small linear model is a lightweight task, and the number of replications is moderate, the overhead dominates and parallel method does not run faster. The work is too small to benefit from parallel methods. If the number of replications is much larger and the model is more computational complex, the parallel computation would work better.

Problem 3 - Large data

(a)

```
Data.frame `df` is 251.8 Mb
```

```
library(lme4)
```

```
Loading required package: Matrix
```

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following object is masked _by_ '.GlobalEnv':
```

```
contains
```

```
The following objects are masked from 'package:stats':
```

```
filter, lag
```

```
The following objects are masked from 'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
countries <- unique(df$country)

results <- list()
coef_forum <- data.frame()

for (c in countries) {
  cat("Fitting model for country:", c, "\n")

  sub <- df %>% filter(country == c)

  sub <- sub %>%
    mutate(
      prior_gpa_z      = scale(prior_gpa),
      forum_posts_z    = scale(forum_posts),
```

```

    quiz_attempts_z = scale(quiz_attempts)
  )

t <- system.time({
  fit <- glmer(
    completed_course ~ prior_gpa_z + forum_posts_z + quiz_attempts_z +
      (1 | device_type),
    data = sub,
    family = binomial()
  )
})

results[[c]] <- list(model = fit, time = t)

coef_forum <- rbind(
  coef_forum,
  data.frame(
    country = c,
    estimate = fixef(fit)["forum_posts_z"],
    time_sec = t["elapsed"]
  )
)
}
}

```

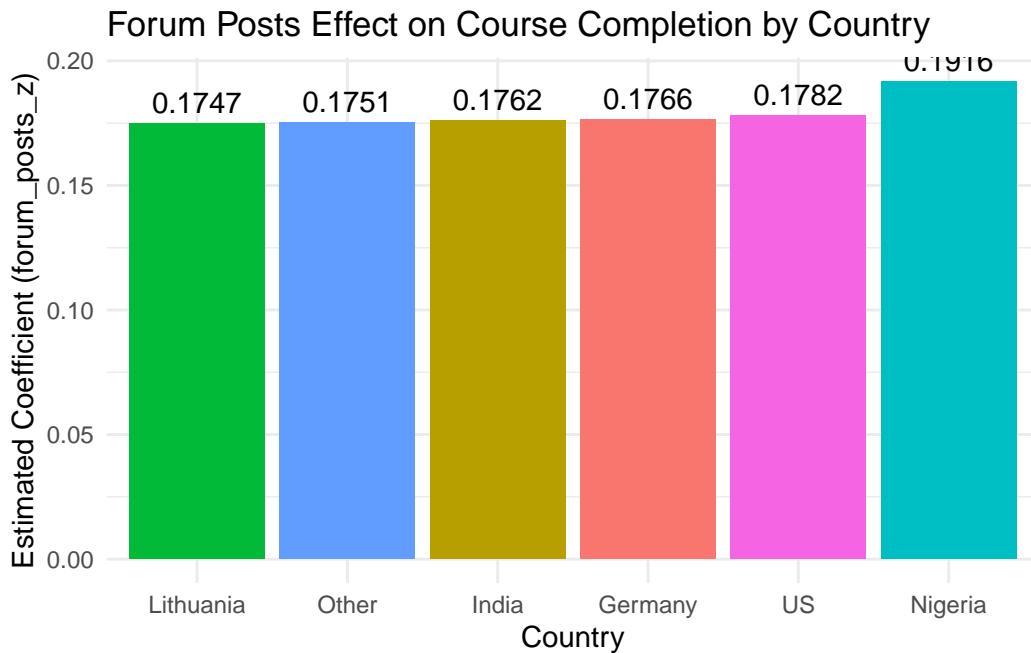
Fitting model for country: US
 Fitting model for country: Other
 Fitting model for country: India
 Fitting model for country: Germany
 Fitting model for country: Lithuania
 Fitting model for country: Nigeria

`coef_forum`

	country	estimate	time_sec
forum_posts_z	US	0.1781708	269.32
forum_posts_z1	Other	0.1751446	509.56
forum_posts_z2	India	0.1761594	206.31
forum_posts_z3	Germany	0.1765869	132.14
forum_posts_z4	Lithuania	0.1747332	3.84
forum_posts_z5	Nigeria	0.1916254	8.57

The running time of the six models are on the table above as time_sec.

```
ggplot(coef_forum, aes(x = reorder(country, estimate),
                        y = estimate,
                        fill = country)) +
  geom_col() +
  geom_text(aes(label = round(estimate, 4)),
            vjust = -0.5,
            size = 4) +
  theme_minimal() +
  labs(
    title = "Forum Posts Effect on Course Completion by Country",
    y = "Estimated Coefficient (forum_posts_z)",
    x = "Country"
  ) +
  theme(legend.position = "none")
```



(b)

```

library(parallel)

detectCores()

[1] 8

total_time <- system.time({

  source("problem_set_6_gen_data.R")

  df2 <- df %>%
    group_by(country) %>%
    mutate(
      prior_gpa_z      = as.numeric(scale(prior_gpa)),
      forum_posts_z    = as.numeric(scale(forum_posts)),
      quiz_attempts_z = as.numeric(scale(quiz_attempts))
    ) %>%
    ungroup()

  fit_country_model_par <- function(c) {
    library(lme4)
    library(dplyr)

    sub <- df2 %>% filter(country == c)

    t <- system.time({
      fit <- glmer(
        completed_course ~ prior_gpa_z + forum_posts_z + quiz_attempts_z +
          (1 | device_type),
        data = sub,
        family = binomial()
      )
    })

    data.frame(
      country = c,
      estimate = fixef(fit)["forum_posts_z"],
      time_sec_fit = t["elapsed"]
    )
  }
})

```

```

countries_list <- unique(df2$country)
num_cores <- detectCores() - 1

cl <- makeCluster(num_cores, type = "PSOCK")

clusterExport(cl, "df2", envir = environment())
clusterEvalQ(cl, {
  library(lme4)
  library(dplyr)
})

parallel_results_list <- parLapply(cl, countries_list, fit_country_model_par)

stopCluster(cl)

optimized_coef_parallel <- bind_rows(parallel_results_list)

})

```

Data.frame `df` is 251.8 Mb

```
cat("Total running time:\n")
```

Total running time:

```
print(total_time)
```

user	system	elapsed
5.47	2.97	639.32

```
print(optimized_coef_parallel)
```

	country	estimate	time_sec_fit
forum_posts_z...1	US	0.1781708	386.16
forum_posts_z...2	Other	0.1751446	622.30
forum_posts_z...3	India	0.1761594	315.91
forum_posts_z...4	Germany	0.1765869	202.67
forum_posts_z...5	Lithuania	0.1747332	5.15
forum_posts_z...6	Nigeria	0.1916254	10.98

The total running time is 800.72 seconds, which is about 13.35 minutes. In (a), the running time is about 40 minutes.

```
coef_a <- coef_forum %>% arrange(country)
coef_b <- optimized_coef_parallel %>% arrange(country)

all.equal(coef_a$estimate, coef_b$estimate)
```

```
[1] TRUE
```

The results of (b) matches those from (a).

Problem 4 - data.table

```
library(data.table)
```

```
Warning: package 'data.table' was built under R version 4.5.2
```

```
Attaching package: 'data.table'
```

```
The following objects are masked from 'package:dplyr':
```

```
between, first, last
```

```
ATP_Matches = fread("https://raw.githubusercontent.com/JeffSackmann/tennis_atp/refs/heads/master/ATP_matches.csv")
```

(a)

```
ATP_Matches[, .N, by = tourney_id][, .N]
```

```
[1] 128
```

128 tournaments are in the *atp_matches_2019* dataset. We assume the tournaments with dates in 2018 are part of tournaments in 2019.

```
library(stringr)

ATP_Matches[str_detect(tourney_name, "Davis Cup")][1:5]
```

	tourney_id	tourney_name	surface						
	<char>	<char>	<char>						
1:	2019-M-DC-2019-FLS-A-M-FRA-JPN-01	Davis Cup Finals RR: FRA vs JPN	Hard						
2:	2019-M-DC-2019-FLS-A-M-FRA-JPN-01	Davis Cup Finals RR: FRA vs JPN	Hard						
3:	2019-M-DC-2019-FLS-A-M-FRA-SRB-01	Davis Cup Finals RR: FRA vs SRB	Hard						
4:	2019-M-DC-2019-FLS-A-M-FRA-SRB-01	Davis Cup Finals RR: FRA vs SRB	Hard						
5:	2019-M-DC-2019-FLS-A-M-SRB-JPN-01	Davis Cup Finals RR: SRB vs JPN	Hard						
	draw_size	tourney_level	tourney_date	match_num	winner_id	winner_seed			
	<int>	<char>	<int>	<int>	<int>	<int>			
1:	4	D	20191119	1	104542	NA			
2:	4	D	20191119	2	106415	NA			
3:	4	D	20191121	1	105936	NA			
4:	4	D	20191121	2	104925	NA			
5:	4	D	20191120	1	105936	NA			
	winner_entry	winner_name	winner_hand	winner_ht	winner_ioc	winner_age			
	<char>	<char>	<char>	<int>	<char>	<num>			
1:		Jo-Wilfried Tsonga	R	188	FRA	34.5			
2:		Yoshihito Nishioka	L	170	JPN	24.1			
3:		Filip Krajinovic	R	185	SRB	27.7			
4:		Novak Djokovic	R	188	SRB	32.4			
5:		Filip Krajinovic	R	185	SRB	27.7			
	loser_id	loser_seed	loser_entry	loser_name	loser_hand	loser_ht			
	<int>	<int>	<char>	<char>	<char>	<int>			
1:	106034	NA		Yasutaka Uchiyama	R	183			
2:	104792	NA		Gael Monfils	R	193			
3:	104542	NA		Jo-Wilfried Tsonga	R	188			
4:	105332	NA		Benoit Paire	R	196			
5:	105216	NA		Yuichi Sugita	R	173			
	loser_ioc	loser_age	score	best_of	round	minutes	w_ace	w_df	w_svpt
	<char>	<num>	<char>	<int>	<char>	<int>	<int>	<int>	<int>
1:	JPN	27.2	6-2 6-1	3	RR	58	4	0	42
2:	FRA	33.2	7-5 6-2	3	RR	66	1	0	51
3:	FRA	34.5	7-5 7-6(5)	3	RR	106	5	1	69
4:	FRA	30.5	6-3 6-3	3	RR	69	3	2	53
5:	JPN	31.1	6-2 6-4	3	RR	68	5	0	47
	w_1stIn	w_1stWon	w_2ndWon	w_SvGms	w_bpSaved	w_bpFaced	l_ace	l_df	l_svpt
	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<int>
1:	27	22	11	8	2	2	3	2	37
2:	34	27	10	10	0	1	6	3	65
3:	49	42	13	12	1	1	18	4	88
4:	33	25	14	9	2	2	6	5	59
5:	31	26	10	9	0	0	3	1	62
	l_1stIn	l_1stWon	l_2ndWon	l_SvGms	l_bpSaved	l_bpFaced	winner_rank		

	<int>						
1:	24	14	5	7	1	5	29
2:	51	31	4	10	2	6	73
3:	56	44	12	12	9	10	40
4:	37	25	8	9	3	6	2
5:	37	21	11	9	6	9	40

	winner_rank_points	loser_rank	loser_rank_points
	<int>	<int>	<int>
1:	1410	81	651
2:	764	10	2530
3:	1148	29	1410
4:	9145	24	1538
5:	1148	104	561

```
tourneys <- ATP_Matches[, tourney_name := str_replace(tourney_name, "Davis.*", "Davis Cup")]
tourneys <- tourneys[, .(tourney_name), by = tourney_name]

nrow(tourneys)
```

[1] 69

There are 69 unique tournaments.

(b)

If we only consider the tournaments with a final round “F”, there are 67 such tournaments. And assume the winner of a tournaments is who wins round “F”.

```
ATP_Matches[round == "F", .N]
```

[1] 67

12 players won more than one tournaments. The most winning players Dominic Thiem and Novak Djokovic both won 5 tournaments.

```
ATP_Matches[round == "F", .N, by = winner_name][N > 1][order(-N)]
```

	winner_name	N
	<char>	<int>
1:	Novak Djokovic	5
2:	Dominic Thiem	5

```

3: Daniil Medvedev      4
4: Roger Federer        4
5: Rafael Nadal         4
6: Alex De Minaur       3
7: Stefanos Tsitsipas   3
8: Jo-Wilfried Tsonga   2
9: Nick Kyrgios          2
10: Cristian Garin       2
11: Benoit Paire          2
12: Matteo Berrettini     2

```

If we consider winners of the round with maximum match number, there are 61 tournaments only having round “RR”.

```
ATP_Matches[order(tourney_id, -match_num), .SD[1], by = tourney_id][, .N, by = round]
```

```

round      N
<char> <int>
1: F      67
2: RR     61

```

```
ATP_Matches[order(tourney_name, -match_num)][, .SD[1], by = tourney_id][, .(wins = .N), by =
```

```

winner_name  wins
<char> <int>
1: Rafael Nadal    9
2: Novak Djokovic  8
3: Alex De Minaur   6
4: Dominic Thiem    5
5: Roger Federer   4
6: Daniil Medvedev  4
7: Cristian Garin   4
8: Denis Shapovalov 4
9: Robin Haase      3
10: Stefanos Tsitsipas 3
11: Nick Kyrgios    2
12: Matteo Berrettini 2
13: Karen Khachanov   2
14: Diego Schwartzman 2
15: Taylor Fritz     2
16: Benoit Paire      2
17: Jo-Wilfried Tsonga 2

```

17 players won more than one tournaments. The most winning player Rafael Nadal won 9 tournaments.

(c)

There is evidence that winners have more aces than losers.

```
library(tidymodels)
```

Warning: package 'tidymodels' was built under R version 4.5.2

-- Attaching packages ----- tidymodels 1.4.1 --

v broom	1.0.10	v rsample	1.3.1
v dials	1.4.2	v tailor	0.1.0
v infer	1.0.9	v tidyverse	1.3.1
v modeldata	1.5.1	v tune	2.0.1
v parsnip	1.4.0	v workflows	1.3.0
v purrr	1.1.0	v workflowsets	1.1.1
v recipes	1.3.1	v yardstick	1.3.2

Warning: package 'dials' was built under R version 4.5.2

Warning: package 'modeldata' was built under R version 4.5.2

Warning: package 'parsnip' was built under R version 4.5.2

Warning: package 'recipes' was built under R version 4.5.2

Warning: package 'rsample' was built under R version 4.5.2

Warning: package 'tailor' was built under R version 4.5.2

Warning: package 'tune' was built under R version 4.5.2

Warning: package 'workflows' was built under R version 4.5.2

Warning: package 'workflowsets' was built under R version 4.5.2

```

Warning: package 'yardstick' was built under R version 4.5.2

-- Conflicts ----- tidymodels_conflicts() --
x data.table::between()  masks dplyr::between()
x .GlobalEnv::contains() masks tidyR::contains(), rsample::contains(), dplyr::contains()
x purrr::discard()      masks scales::discard()
x tidyR::expand()       masks Matrix::expand()
x dplyr::filter()       masks stats::filter()
x data.table::first()   masks dplyr::first()
x recipes::fixed()     masks stringr::fixed()
x dplyr::lag()          masks stats::lag()
x data.table::last()    masks dplyr::last()
x tidyR::pack()         masks Matrix::pack()
x rsample::populate()  masks Rcpp::populate()
x recipes::step()       masks stats::step()
x purrr::transpose()   masks data.table::transpose()
x tidyR::unpack()       masks Matrix::unpack()
x recipes::update()    masks Matrix::update(), stats::update()

ATP_Matches[, mean(w_ace > l_ace, na.rm = TRUE)]

```

```
[1] 0.5727543
```

```

prop_test(ATP_Matches[, .(win_more = w_ace > l_ace)], win_more ~ NULL, p = 0.5)

# A tibble: 1 x 4
  statistic chisq_df  p_value alternative
    <dbl>     <int>    <dbl> <chr>
1      56.7       1 4.95e-14 two.sided

```

We define `win_more = (w_ace > l_ace)` and conduct a one-sample proportion test with $H_0: P(\text{win_more}) = 0.5$, $H_1: P(\text{win_more}) \neq 0.5$.

The sample proportion of matches in which winners record more aces is: `mean(win_more) = 0.57 > 0.5`

Since $p\text{-value} = 4.95e-14 \ll 0.0001$, $t\text{-statistic} = 56.7487$, H_0 is rejected, and winners are significantly more likely than losers to record more aces.

(d)

The player (at least 5 matches) with highest win-rate is *Rafael Nadal*.

```

ATP_Matches[
  , melt(.SD,
    measure.vars = c("winner_name", "loser_name"),
    variable.name = "outcome",
    value.name   = "player_name")
][
  , .(
    total = .N,
    wins  = sum(outcome == "winner_name")
  ),
  by = player_name
][
  total >= 5
][
  , win_rate := wins / total
][
  win_rate == max(win_rate)
]

player_name total  wins  win_rate
<char> <int> <int>      <num>
1: Rafael Nadal    69     60 0.8695652

```